

Search (state space: how to rep. / f(two))

A*: frontier rep. as priority Q: sorted by lowest est. Cost

- $f(n) = g(n) + h(n)$; $g(n)$: dist. traveled already; $h(n)$: est. cost to goal
- heuristic informs about knowledge on world

Admissibility: $\forall n, 0 \leq h(n) \leq h^*(n)$; $h^*(n)$ is True cost to goal

- Under estimates Cost to goal. "optimistic"

Consistency: $\forall A, C, h(A) - h(C) \leq \text{Cost}(A, C)$

- Under Estimates Cost to next node. triangle inequality.
- Stronger than Admissibility: Consistency \Rightarrow Admissibility

GAMES

Zero-sum: Nodes take turns (Max & Min if optimal)

(Minimax): $b^m = O(n)$

Alpha-beta pruning: $O(b) = b^{m/2} \rightarrow$ we can prune if there is

- if node is random, we don't need to look at it. No chance it meets Criterion
- other parents

Expectimax: random node can cause \neq optimal

- Be carefull about Pruning.

$E \rightarrow$

LLOGIC use Knowledge base (from observations) to derive stories

Logical Eq

$$\begin{aligned} \alpha \wedge \beta &\equiv \beta \wedge \alpha \\ \alpha \vee \beta &\equiv \beta \vee \alpha \\ (\alpha \wedge \beta) \wedge \gamma &\equiv \alpha \wedge (\beta \wedge \gamma) \\ (\alpha \vee \beta) \vee \gamma &\equiv \alpha \vee (\beta \vee \gamma) \\ \neg(\neg \alpha) &\equiv \alpha \\ \alpha \Rightarrow \beta &\equiv \neg \beta \Rightarrow \neg \alpha \\ \alpha \Rightarrow \beta &\equiv \neg \alpha \vee \beta \\ \alpha \Leftrightarrow \beta &\equiv (\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha) \end{aligned}$$

$$\begin{aligned} \neg(\alpha \wedge \beta) &\equiv \neg \alpha \vee \neg \beta \\ \neg(\alpha \vee \beta) &\equiv \neg \alpha \wedge \neg \beta \\ \alpha \wedge (\beta \vee \gamma) &\equiv (\alpha \wedge \beta) \vee (\alpha \wedge \gamma) \\ \alpha \vee (\beta \wedge \gamma) &\equiv (\alpha \vee \beta) \wedge (\alpha \vee \gamma) \end{aligned}$$

Model: T, f assignments for any given problem

Valid: True in all worlds / models

Satisfiable: True in one world

Unsatisfiable: Never true ($\neg A \wedge A$)

entailment (\models): $A \models B$: If in All worlds A is true, B is true

- Direct Proof**: $A \models B$: If $A \Rightarrow B$ is Valid
- Contradiction**: $A \models B$ iff $A \wedge \neg B$ is unsatisfiable

CNF: Conjunctive normal form: only in: $\wedge, \neg, \vee, (,)$

Model Checking: $\models B \vdash a$

- Enumeration / Truth table**
- Table over true / false and find out

SAT vs DPLL: Checks Satisifiability (find 1 world is B true) \rightarrow to do: find a world (variable assignment) to make state true

Theorem Proving

- $\neg \neg A \Rightarrow B$, if we know A, we know B
- $\neg \neg A \wedge B$, if we infer A, we infer B
- if we have $A \wedge B$, we infer $A \wedge B$
- Forward chaining**: implication chain

Guaranteed optimal path

DPS: X

BFS: if Cost = 1

A*: if Consistency Criterion

equal: if Cost = 1, BFS = UCS

Monte Carlo tree Search

- randomly
- uses random decisions
- learns certain faster
- selection, exploitation, Sim, Back Prop.

Completeness

BFS, A^* (but max)

Tree search V.s. Graph search

- Admissible heuristic opt. for tree $\not\approx$ graph
- generalized

Tree search = graph search + [reached set]

A is Necessary & Sufficient for B

$$\Rightarrow \Leftarrow \rightarrow \Leftarrow \Leftrightarrow$$

using rule of proportionality (\propto)

$$\begin{aligned} &\rightarrow \neg x \models \text{find } P(+x|b), P(-x|b) \\ &\rightarrow \frac{P(0|+x)P(b)}{P(0|b)} \leftarrow \frac{P(b|-\neg x)}{P(b)} \\ &\left(\frac{P(+x)}{P(b)} \right) \quad \text{OR } P(+x, b) = P(+x)P(b|x) \\ &\quad \rightarrow \text{find answer then normalize} \\ &\quad \frac{P(+x, b)}{P(+x, b) + P(-x, b)} \leftarrow \text{Linearization} \quad \frac{(-x, b)}{P(+x, b) + P(-x, b)} \\ &\quad \text{Simpler State Condition:} \end{aligned}$$

$F^{T+1} \leftarrow \text{Action clauses}(F^T) \vee (F^+ \wedge \text{Achieves} \text{Loc}(F^+))$

Model: $\{S : T, N = F\}$

Symbol: $\{L, M, N, P, Q, R, S\}$

Atoms: $(\neg L \vee \neg S) \wedge (M \vee Q \vee N) \wedge (L \vee M) \wedge (L \vee G)$

Atoms: $(\neg L \vee \neg P) \wedge (Q \vee P \vee N) \wedge (\neg R \vee \neg L) \wedge S$

Model: $\{S : T, N = F\}$

Symbol: $\{L, M, N, P, Q, R\}$

Atoms: $(\neg L \vee \neg T) \wedge (M \vee Q \vee N) \wedge (L \vee M) \wedge (L \vee Q) \wedge (\neg L \vee \neg P)$

Atoms: $\neg (R \vee P \vee N) \wedge (\neg R \vee \neg L)$

Model: $\{S : T, N = F\}$

Atoms: $(M \vee Q \vee F) \wedge (L \vee M) \wedge (L \vee Q) \wedge (\neg L \vee \neg P)$

Atoms: $\neg (R \vee P \vee F) \wedge (\neg R \vee \neg L)$

Choose: $M = \text{True}$, then $M = \text{False}$

If you get $(Q) \wedge (\neg Q)$ or some unsatisfiable clause, Unsatisfiable, else: Satisfiable

$A \rightarrow B$, $A \rightarrow C$, $B \wedge C \rightarrow D$, $D \wedge E \rightarrow Q$, $A \wedge D \rightarrow Q$
 Inference: [A] \vdash [B, C] \vdash [D] \vdash [Q], $A \wedge D \rightarrow Q$
 $B \wedge C \rightarrow D$, $D \wedge E \rightarrow Q$, $A \wedge D \rightarrow Q$
 Inference: [A, B, C] \vdash [C] \vdash [D] \vdash [Q]
 $D \wedge E \rightarrow Q$, $A \wedge D \rightarrow Q$
 Inference: [A, B, C, D] ...

Probability

$$P(a|b) = \frac{P(b|a)P(a)}{P(b)}$$

$$\text{Independence: } P(x|y) = P(x); P(y|x) = P(y)$$

Unitary Constraint uses 1 var...
Binary Constraint uses 2 var ...

Conditional Independence: X is cond. indep of Y given Z : $P(X|Y, Z) = P(X|Z)$
 $P(X, Y|Z) = P(X|Z)P(Y|Z)$

BAYES NETS

Variable Elimination: join all facts of X
 \rightarrow sum out X .

Factor: any unnormalized prob

Inference by Enumeration

Visit all out: CPT: $P(B), P(E)$, $P(A|B, E)$, $P(V|A)$, $P(M|A)$

$P(G) = \sum_{n \in N} \sum_{d \in D} \sum_{i \in I} \sum_{a \in A} P(M=n)P(C=d)$

$P(G) = \sum_{i \in I} \sum_{d \in D} \sum_{a \in A} P(M=i)P(C=d)$

$P(G) = \sum_{a \in A} P(M=a)P(C=a)$

Prior Sampling (bigness)

Open samples,

\rightarrow add samples in consistent w/ evidence at end

\rightarrow calc. prob.

Rejection Sampling

Open samples,

\rightarrow add samples in consistent w/ evidence at end

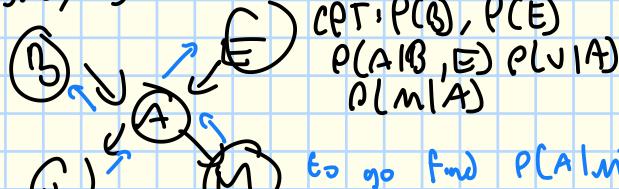
as we generate

\rightarrow calc. prob.

Likelyhood Weighting

Fix evidence var. to a 200k-learn (height or weight or multipl. of dist., but vars)

\rightarrow multiply out dist., but vars by prob



To go fwd: $P(A|M)$ use Bayes rule: $\frac{P(M|A)P(A)}{P(M)}$

$$P(C=c|B=b) = P(c|a,b) + P(\neg c|a,b) + P(c|a,\neg b) + P(\neg c|a,\neg b)$$

$\boxed{A \perp\!\!\!\perp B | C}$ if $P(B|A) = P(B)$
 can still be true

Independence

\rightarrow given parents, node is indep. to all other children nodes

\rightarrow given Markov blanket, each node is cond. indep. (parents, children, children's other parents)

Maximising $P(Q)$ given constraints

\rightarrow Make polynomial of new Markov blanket given constraints

\rightarrow solve derivative of $P(Q)$ wrt local max. along $\frac{\partial f(Q)}{\partial Q} = 0$

Let $q = P(+n)$ represent the probability, and then $P(+d) = (1 - q)$. We then know that $P(-n) = (1 - q)$ and $P(-d) = 1 - (1 - q) = q$. Then we have:

$$P(+u) = 0.8q(1 - q) + 0.3q^2 + 0.6(1 - q)^2 + 0.7(1 - q)q$$

$$P(+u) = -0.6q^2 + 0.3q + 0.6$$

To optimize this expression, we take the first derivative and set it to zero. Since it is a quadratic with negative leading coefficient, this point would be a maximum.

Markov Blanket: each node is inter to all nodes given its Markov Blanket. \rightarrow Parents, Children, Children's other parents

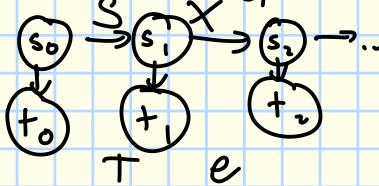
Gibbs Sampling

\rightarrow Fix ex. variables. (all other vars are true)

\rightarrow open new stroke by looping through new

variables given samples.

HMMs hidden(S); observe(T) \Rightarrow Viterbi Search Path

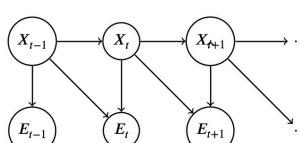


\rightarrow Most likely sequence S_0, S_1, S_2, \dots to produce observe chunk = argmax $y \cdot P(y) \prod_i P(s_i, y_i)$
 \rightarrow if $(-\log(P))$, go with smaller prob; by transition
 \rightarrow if $(\log(P))$, we could maximize our states

Bayes classifier

$$\int_{\text{new features}} p(\text{class} | \text{new features})$$

(c) [4 pts] After switching to a low-cost sensor, the observations are noisier. One way to address the increased uncertainty in the observations is to condition the observation on more than one state. Complete the forward algorithm updates for an HMM in which the observation E_t is conditioned on both the current state X_t and the previous state X_{t-1} .



$$P(x_t | e_{1:t}) = \sum_{x_{t-1}} P(e_t | x_{t-1}) P(x_t | x_{t-1}) P(x_{t-1} | e_{1:t-1})$$

$$T_1 \perp\!\!\!\perp S_0 | S_1$$

$$P(+, | S_0) = P(+, | S_0, S_1)$$

$$\text{Belief distribution}$$

$$B(S_n) = P(S_n | t_0, +, \dots, t_{n-1})$$

$$B'(S_n) = P(S_n | t_0, +, \dots, t_{n-1})$$

$$\text{Forwarding:}$$

$$g(w_{in}) = P(w_{i+1} | w_i) g(w_i)$$

$$g(w_{in}) \propto P(f_{i+1} | w_{in}) b(w_{i+1})$$

$$g(w_{in}) \propto P(f_{i+1} | w_{in}) \sum_w P(w_{in} | w) b(w)$$

Expected utilities: $\text{VE}^*(e) \geq \text{VPI}(E|e) \geq 0$ Bellman eq.

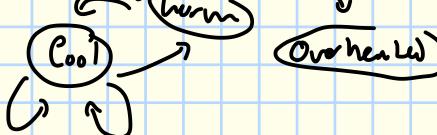
(M)ax(E) \times $\text{optimal}(e)$ utility = Max util given choice $\rightarrow \text{VPI}(E_j, E_k, e) = \text{VPI}(E_j|e) + \text{VPI}(E_k|e)$

(V)alue of (P)erf or (I)nfo. $\text{VPI}(E_j, E_k|e) = \text{VPI}(E_j|e) + \text{VPI}(E_k|e, E_j)$

$\text{VPI}(E, e) = \text{MEU}(e, E) - \text{MEU}(e)$ ~ can be negative $\rightarrow \text{VPI}(E_k, e) + \text{VPI}(E_j|e, E_k)$

$\cdot MDP$ $\rightarrow \gamma, \gamma^2, \gamma^3 \rightarrow$ future stakes matter less:

States, Actions, discount factor, transition function, reward function



transition functn. $T(S, a, S') = \text{prob of going to } S'$
state, action, new state
reward $R(S, a, S') = \text{reward}$

$T(S, a, S') = P(S'|S, a)$

Policy $\pi : S \rightarrow A$, for each state \rightarrow action

State optimal value $(U^*(S))$: Best expected Utility at optimal agent @ S

Optimal Q-value $(Q^*(S, a))$: Expected Utility of optimal agent @ S taking an action a in state S

- Transition function $T(S, a, S')$: Probability of taking action a @ S resulting in S'

- Reward functn $R(S, a, S')$: Pos/neg reward for each stake

- terminal state S_T

- start state S_0

$\frac{\text{Reward}}{1-\gamma}$

$Q^*(S, a) = \sum_s T(S, a, s)[R(S, a, s) + \gamma U^*(s)]$

$U^*(S) = \max_a Q^*(S, a)$

$R = \text{reward}$

$Q^* = \text{optimal Q-state}$

$U^* = \text{optimal value @ S}$

Value Iteration \rightarrow Opt alg. to compute MDP

- 1) Vs GS, set $U_0(S) = 0$
- 2) Repeat till converge. $\mathcal{O}(S^2 \cdot A)$

\rightarrow if $\gamma=1$, doesn't converge. If $\gamma < 1$ converge

Vs GS, $U_{k+1}(S) \leftarrow \max_S \sum S T(S, a, S') [R(S, a, S') + \gamma U_k(S)]$

Policy Iteration \rightarrow More efficient metric to

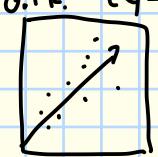
- 1) def initial policy π_0 converge to opt. policy
- 2) use policy eval. for curr. policy $\mathcal{O}(S^3)$
- 3) use policy improvement

$U^\pi(S) = \sum_S T(S, \pi(S), S) [R(S, \pi(S), S) + \gamma U^\pi(S)]$

$\pi_{i+1}(S) = \arg\max_a \sum_a T(S, a, S) [R(S, a, S) + \gamma U^\pi(S)]$

Linear Regression \rightarrow Loss of our funct. square diff. $(y - h(x))^2$

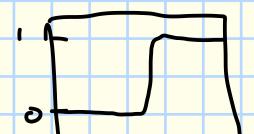
larger squares $\Rightarrow \hat{w} = (x^T x)^{-1} x^T y$



Log regression \rightarrow classifier for continuous vars.

$\rightarrow h_w(x) = \frac{1}{1 + e^{-w^T x}}$

\rightarrow gradients $p(\text{class} | \text{evidence})$



Naive Bayes \rightarrow Model features as Bayes, assume each feature indep from each other

$\hookrightarrow \text{prediction}(F) = \arg\max_i p(y_i | F) \prod_j p(F_j | f_j, Y=y_j)$, normalize

f_j = features, y_i = prediction, features are cond. indep (some prob known indep)

Maximum Likelihood Estimation (MLE) \rightarrow given i.i.d., Method to learn probability, count n

\rightarrow Likelihood: $L(\theta) = \prod_{i=1}^n p_\theta(x_i)$ take gradient

$\frac{\partial}{\partial \theta} L(\theta) = 0$ to get max, log likelihood

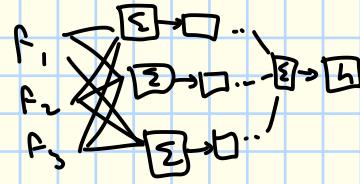
$\log L(\theta) \in$

\rightarrow Laplace Smoothing: mitigate overfitting, assume seen K extra each of each outcome

$P_{\text{Laplace}}(x) = \frac{\text{Count}(x) + k}{N + k|x|}$

$|X| = \text{set of pos. values}$

Neural Networks



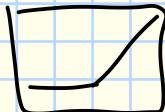
Multilayered Perception: maps down to higher dimension then classify
Universal func approximator: two layer NN w/sufficient neurons can approx
any cont. func can use indicator function or sgn threshold, softmax func to classify

$$\text{log likely hood} = \text{log l}(\omega) = \prod_{i=1}^n \log P(y_i | f(x_i), \omega)$$

Activation functions

$$\text{Sigmoid: } \sigma(x) = \frac{1}{1+e^{-x}}$$

$$\text{ReLU: } f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$



Backprop.
determine gradient of output w/ respect to each of input

1) Forward Pass: Compute Values through Computation graphs

2) Backwards Pass: Compute gradients taking advantage of chain rule

Reinforcement Learning

Episode: Collection of Samples which are (S, A, S', r) tuples

Model Based Learning: estimates transitions, reward functions w/ samples before using policy / value if

→ Count times arrived in state and normalized counts, Law of large numbers will converge on opt.

→ Model-free learning: estimate Q-values directly w/o constructing model

↳ Direct Evaluation: fix policy π and have agent experience, can compute by averages

↳ Temporal diff learning: Learning from every experience

$$\text{Sample} = R(S, \pi(S), S') + \gamma V^\pi(S)$$

$$V^\pi(S) \leftarrow (1-\alpha) V^\pi(S) + \alpha \cdot \text{Sample}$$

→ Active Reinforcement Learning: learning agent agent can use feed back to iteratively update policy

↳ Q-learning: bypass model by directly learning Q-values

Q-Value Iteration

$$Q_{\text{new}}(S, a) \leftarrow \sum_{S'} T(S, a, S) [R(S, a, S) + \gamma \max_a Q(S', a)]$$

Q-value Samples:

$$\text{Sample} = R(S, a, S') + \gamma \max_a Q(S', a)$$

$$Q(S, a) \leftarrow (1-\alpha) Q(S, a) + \alpha \cdot \text{Sample}$$

Feature Based Representation: allows model to generalize learning experiences, store as linear feature

$$\text{Weight update: } w_i \leftarrow w_i - \alpha \cdot f_i(S, a) \cdot \frac{\partial Q_w}{\partial w_i} \quad \left. \right\} \text{negative for grad. decent}$$

Exploration vs Exploitation

$$\frac{\partial Q_w}{\partial w_i} = [(R(S, a) + \gamma \max_a Q(S', a))] - Q(S, a)$$

ε-greedy: $0 \leq \epsilon \leq 1$, act randomly w/ prob ϵ , should be lowered over time

Exploration functions:

$$(Q(S, a) = (1-\alpha)Q(S, a) + \alpha[R(S, a, S') + \gamma \max_a Q(S', a)])$$

modification update

$$P(S, a) = Q(S, a) + \frac{1}{N(S, a)} \quad \leftarrow \text{Predetermined}$$

of times Q is visited

Replay: metric to measure model, difference between total reward for option and reward in our model

